SS09, 2010
Lab Assignment 5: Code Optimization

Introduction

This assignment requires you to tune the performance of some simple C functions. There are two parts to the assignment. The first involves code where performance is constrained by arithmetic operations. In the second part, performance is constrained by memory operations.

Logistics

Everyone should work individually on this assignment. The only thing handed in will be your version of the two files **poly.c** and **rowcol.c**, which you must submit via our SVN server (10.132.143.1) like in Lab0.

You can retrieve the necessary files for this assignment from the server. There are a number of files. The only two that you will edit are **poly.c** and **rowcol.c**

Part I: Polynomial evaluation

A polynomial of degree d consists of d+1 coefficients $a_0, a_1, \ldots, a_d$. Given some value **X**, we evaluate the polynomial by computing:

$$P(x) = a_0 + a_1 x + a_2 x^2 + \cdots + a_d x^d$$

For integer coefficients and values of **X**, this can be performed by the

following code:

```
int poly_eval(int *a, int degree, int x)
{
    int result = 0;
    int i;
    int xpwr = 1; /* Successive powers of x */
    for (i = 0; i <= degree; i++) {
        result += a[i]*xpwr;
        xpwr *= x;
    }
    return result;
}
```

In this code, we achieve linear complexity by computing successive powers of

x as we evaluate the coefficients in order. We would then expect the total

number of cycles to evaluate a polynomial of degree to be of the form:

$$C(d) = \alpha d + \beta$$

We will measure the performance of the evaluation function in two ways:

CPE: The number of clock cycles per element, that is, the term in the

equation shown above. This measure is useful for analyzing how the

performance grows for high degree polynomials.

C(10): The total number of clock cycles to evaluate a degree-10 polynomial.

This measure is useful for analyzing the performance of the function on a

polynomial of more typical degree. Our original implementation has a CPE of

11.00, and a C(10) of **160**. (Tested on machines in room 303 of software

building).

Your job is to create versions of the functions that minimize the two

measures. You will find that the two measures call for somewhat different

optimizations.

One technique you might find useful is known as <u>Horner's Rule</u>. By exploiting the associative, commutative, and distributive laws, we can write the polynomial of Equation 1 as:

$$P(x) = a_0 + (a_1 + (a_2 + (\cdots + (a_{d-1} + a_d x)x \cdots)x)x)x$$

This reduces the number of multiplications to around 1/2 of what is required in the reference code shown above. You will find that a direct implementation of this rule actually slows down the function, but you can speed things up with the usual sorts of optimizations, such as loop unrolling.


Figure 1: Computation of Performance Score. Your score depends on how much speedup your implementation gives over the reference version, compared to that of our best implementation

You will find it best to try a number of different implementations, with different degrees of unrolling, different associations of operations, etc. Our testing code provides a convenient way to test (for correctness) and measure (for performance) your many variants. That is, you can implement any number of functions with the calling convention shown in the code for poly_eval within the file **poly.c**. At the end of this file you will find a table, in the form of an initialized array, where you can put entries for different implementations of the function. You will find this is a useful way to keep a

record of the different solutions you create. You should arrange your entries in the table so that the function with lowest CPE comes first, and the function with lowest C(10) comes second. The initial version of the table lists function poly_eval for these two entries.

To test and measure your implementations, use the command "**make poly_test**" to compile the test program and then run this program with the command "**./poly_test**". It will test each function in your table for correctness and then compute its CPE and C(10) values.

Your grade for this part will be based on two measures: performance and code quality. Performance will be measured by how much speedup your solution gives over the original solution. That is, if your program requires C_{meas} cycles (either CPE or C(10)), then your speedup is $S_{meas} = C_{ref} / C_{meas}$, where C_{ref} is the number of cycles required by the reference implementation. Similarly, we have made our best attempt at optimizing the function, giving a cycle time of C_{best} and a speedup of $S_{best} = C_{ref} / C_{best}$. As shown in Figure 1, you will be given a score on a 100 point scale according to how well your speedup does relative to S_{best}. To get a nonzero score, you must achieve at least **10%** of the speedup that we did (Slow). To get a score of **100**, you must achieve **90%** of the speedup we did (S_{full}). If you can match us, or even outperform us, you will get extra credit up to a maximum score of **120** (up to S_{high}). For your convenience, the scores for the two functions are printed when you run poly_test.

	Polynomial Evaluation		Matrix Summation	
	CPE	C(10)	Column CME	Row/Col.CME
C_{ref}	11	160	21	27
C_{best}	2.45	135	3.10	5.50
S_{best}	4.49	1.19	6.77	4.91
C_{low}	8.15	157.09	13.31	19.41
C_{full}	2.66	137.14	3.39	5.98
C_{high}	2.27	132.92	2.86	5.09

Figure 2: Performance Standards. C_{ref} indicates the performance (in clock cycles) of the reference implementation. C_{best} indicates the performance of the best implementation we have found. C_{low} indicates the performance required to get any performance credit (score = 20). C_{full} indicates the performance required to get full performance credit (score = 100). C_{high} indicates the performance required to get maximum performance credit (score = 120).

Figure 2 shows the actual cycle counts required to meet the thresholds low (score = 20), full (score = 100), and high (score = 120).

The code quality will be determined by looking at your code. Factors we will consider include the quality of the comments, choices of variable names, etc. Good code style is very important when you write code that does a lot of tricks for optimization. Include in your comments an explanation of the transformations you made and why they help improve the program performance. We will only look at the code for your two solution functions (best CPE and best C(10)).

Part II: Matrix Code

Consider the following definition of a square integer matrix:

```
#define N 512
/* N x N matrix */
typedef int matrix_t[N][N];
/* Vector of length N */
typedef int vector_t[N];
```

That is, data type *matrix_t* denotes a 512 * 512 matrix of integers, while data type *vector_t* denotes a vector of 512 integers. Given a matrix M, we wish to

implement two different functions:

Column Sum: Compute a vector, where entry *i* is the sum of all elements in column of M.

Row/Column Sum: Compute two vectors. In one, entry *i* is the sum of all elements in row *i* of while in the second, entry *i* is the sum of all elements in column of M.

The following is our reference implementation for the column summation:

```
void c_sum(matrix_t M, vector_t rowsum, vector_t colsum)
{
    int i,j;
    for (j = 0; j < N; j++) {
        colsum[j] = 0;
        for (i = 0; i < N; i++)
            colsum[j] += M[i][j];
    }
}
```

Observe that the second argument to this function, <u>rowsum</u>, does not get used at all. We use this calling convention to enable the same testing code to be used with both types of summation functions.

The following is our reference implementation for the row/column summation:

```
void rc_sum(matrix_t M, vector_t rowsum, vector_t colsum)
{
    int i,j;

    for(i=0;i<N;i++){
        rowsum[i] = colsum[i] = 0;
        for(j= 0;j < N;j++){
            rowsum[i] += M[i][j];
            colsum[i] += M[j][i];
        }
    }
}
```

We will measure the performance of the two functions in terms of the

number of clock cycles per *matrixelement* (CME). If the function requires C total cycles, then the performance is computed as CME = C / (512*512), We use this form of measurement rather than CPE, since the underlying assumption of CPE is that performance scales in a simple, linear way with the number of elements. When code encounters cache effects, the performance varies nonlinearly.

Our reference implementation of *c_sum* requires around 21 cycles per element, while while *rc_sum* requires around 27(Tested on machines in room 303 of software building).

As in Part I, you can implement as many different column or row/column summation functions as you like in the file **rowcol.c**. You then enter the different function names and their descriptions in the table at the end of the file. Use the enumerated types COL and ROWCOL to indicate the type of summation function. Arrange the table so that your best column summation function comes first and your best row/column summation function comes second.

To test and measure your implementations, use the command "**make rowcol_test**" to compile the test program and then run this program with the command "**./rowcol_test**". It will test each function in your table for correctness and then compute its CME value.

As in Part I, your grade for each part will be determined by a combination of code quality and performance. Performance for a function will be measured

by your speedup relative to that of our best implementation. The cycle thresholds are shown in Figure 2. Your performance scores are printed when you run "./**rowcol_test**".

Additional Notes

All of the code you write must be in standard, ANSI C. You may not use any embedded assembly or other features specific to GCC.

You must develop and test your program on any IA32 machine.

You will find that the measured performance deviates from one run to the next. This is especially likely to happen with the memory-intensive matrix code when running on a heavily loaded system. Try to work at a time when the machines are lightly loaded to get accurate measurements.

Special Note:

The scores printed on the screen when you do the lab depend on the ref and best CPE or C(10) or CME on our testing server. So when measuring the performance on your own machines, the scores may be somehow low or high. Just focus on how much the CPE or C(10) or CME is improved. The more the performance is improved the higher your score is. TA will give the final score by testing your codes on machines in room 303 of software building.